

JS Engine Security in 2025

New Bugs, New Defenses

Samuel Groß, Google Project Zero



CVE-2019-17026: Zero-Day Vulnerability in Mozilla Firefox Exploited in Targeted Attacks

The Million Dollar Dissident NSO Group's iPhone Zero-Days used against a UAE Human Rights Defender

By Bill Marczak and John Scott-Railton

August 24, 2016

Security

Google patches type confusion zero-day in Chrome's V8 engine

The bug was discovered by its government-backed attacks focused research group.

Responding to Firefox 0-days in the wild

By Philip Martin

Corporate, August 8, 2019, 7 min read time

Google Patches Chrome Zero-Day CVE-2025-10585 as Active V8 Exploit Threatens Millions

Sep 18, 2025

Ravie Lakshmanan

Vulnerability / Browser Security

CVE-2019-17026: Zero-Day Vulnerability in Mozilla Firefox Exploited in Targeted Attacks

> 10 Years of JavaScript Engine Exploitation

The Million Dollar Bug: NSO Group's iPhone Zero-Days used against Human Rights Defender

By Bill Marczak and John Scott-Railton August 24, 2016

Google patches type confusion zero-day in Chrome's V8 engine

The bug was discovered by its government-backed attacks focused research group.

Responding to Firefox bugs in the wild

By Philip Martin Corporate, August 8, 2019, 7 min read time

Google Patches Chrome Zero-Day CVE-2025-10585 as Active V8 Exploit Threatens Millions

Outline

1. Why JavaScript engine's are **hard** to secure
2. Brief history of JavaScript engine exploitation
3. Overview of current and future defenses
4. JS engine vulnerability research in 2025

“Classic” Memory Safety Bugs

```
int array[100];  
int get(int i) {  
    if (i >= 100) return 0;  
    return array[i];  
}
```

“Classic” Memory Safety Bugs

```
int array[100];  
int get(int i) {  
    if (i >= 100) return 0;  
    return array[i];  
}
```

```
get(-1);
```



“Classic” Memory Safety Bugs - Compiler Mitigation

```
int array[100];  
int get(int i) {  
    if (i >= 100) return 0;  
    __bounds_check__(size_t{i}, 100);  
    return array[i];  
}
```



“Classic” Memory Safety Bugs - Compiler Mitigation

```
int array[100];  
int get(size_t i) {  
    if (i >= 100) return 0;  
    return array[i];  
}
```

“Classic” Memory Safety Bugs - Compiler Mitigation

```
int array[100];  
  
int get(size_t i) {  
    if (i >= 100) return 0;  
    __bounds_check__(size_t{i}, 100);  
    return array[i];  
}
```

“Classic” Memory Safety Bugs - Compiler Mitigation

```
int array[100];  
  
int get(size_t i) {  
    if (i >= 100) return 0;  
    --bounds_check__(size_t(i), 100);  
    return array[i];  
}
```



Redundant bounds-check will be removed by compiler
during code optimization ⇒ Low overhead 

JavaScript Engine Bug

```
let array = new Array(100);
function get(i, m) {
    if (i < 0 ||  
        i >= array.length) return;  
    return Number(m) * array[i];
}
```

JavaScript Engine Bug

```
let a = new Array(100);
function get(i, m) {
    if (i < 0 ||  
        i >= array.length) return;  
    return Number(m) * array[i];
}
```

JavaScript Engine Bug

```
let a = new Array(100);
function get(i, m) {
    if (i < 0 ||  
        i >= array.length) return;  
    return Number(m) * array[i];
}
```



```
FUNC get(i, m):  
    SPECULATE typeof(i) == Smi  
    IF i < 0 RETURN  
    IF i >= array.length RETURN  
    M = ToNumber(m)  
    _bounds_check_(i, array.length)  
    A = array[i]  
    RETURN M * A
```

JavaScript Engine Bug

```
let a = new Array(100);
function get(i, m) {
    if (i < 0 ||  
        i >= array.length) return;  
    return Number(m) * array[i];
}
```

```
FUNC get(i, m):  
SPECULATE typeof(i) == Smi  
IF i < 0 RETURN  
IF i >= array.length RETURN  
M = ToNumber(m)  
_bounds_check_(i, array.length)  
A = array[i]  
RETURN M * A
```

```
FUNC get(i, m):  
SPECULATE typeof(i) == Smi  
IF i < 0 RETURN  
IF i >= array.length RETURN  
M = ToNumber(m)  
_bounds_check_(i, array.length)  
A = array[i]  
RETURN M * A
```

JavaScript Engine Bug

```
let a = new Array(100);
function get(i, m) {
    if (i < 0 ||  
        i >= array.length) return;  
    return Number(m) * array[i];
}
```

```
FUNC get(i, m):  
SPECULATE typeof(i) == Smi  
IF i < 0 RETURN  
IF i >= array.length RETURN  
M = ToNumber(m)  
_bounds_check_(i, array.length)  
A = array[i]  
RETURN M * A
```

```
FUNC get(i, m):  
SPECULATE typeof(i) == Smi  
IF i < 0 RETURN  
IF i >= array.length RETURN  
M = ToNumber(m)  
_bounds_check_(i, array.length)  
A = array[i]  
RETURN M * A
```

JavaScript Engine Bug

```
let array = new Array(100);
function get(i, m) {
    if (i < 0 ||
        i >= array.length) return;
    return Number(m) * array[i];
}
let evil = { valueOf() { array.length = 0; } };
get(42, evil);
```



JavaScript Engine Bug


```
FUNC get(i, m):  
    SPECULATE typeof(i) == Smi  
    IF i < 0 RETURN  
    IF i >= array.length RETURN  
    M = ToNumber(m)  
    __bounds_check__(i, array.length)  
    A = array[i]  
    RETURN M * A
```

Redundant bounds-check will be removed by compiler during code optimization...

Except that it's not redundant in this case 😬


Why JavaScript Engine Security is hard

“Typical” Application




Why JavaScript Engine Security is hard


“Typical” Application



JavaScript Engine




Why JavaScript Engine Security is hard



Brief History of JS Engine Security

V8 Architecture




Brief History of JS Engine Bugs (Subjective)

- Phase 1 (<= ~2017): “Classic” runtime bugs
 - Bugs mostly in the runtime and builtin functions
 - Fairly local and somewhat shallow bugs
- Phase 2 (>= ~2017): Optimization bugs
 - Bugs deeper in the execution pipeline (e.g. JIT)
 - Some of the most complex bugs in software security?
- Phase 3 (>= ~2023): Wasm enters the picture
 - Major leap in complexity with WasmGC proposal
 - Brought complexity close to that of the JS pipeline
- Today: Mostly a mix of 2 and 3 (for the major engines)

Phase 1: Runtime Bugs

- By now “classic” JS engine bugs
- Typically bug pattern: unexpected callback in runtime functions which violates previous assumptions
 - But also: integer overflows, etc.
- Bugs often local to a single function
- Have *mostly* disappeared by now in the major JS engines




```
var a = [];
for (var i = 0; i < 100; i++) {
  a.push(i + 0.123);
}
var evil = {
  valueOf: function() {
    a.length = 0;
    return 10;
  }
};
var b = a.slice(0, evil);
```



Phase 2: Optimization Bugs


- Started with fairly simple (in retrospect!) JIT bugs
 - E.g. invalid bounds-check elimination
- Gradually became more and more complex, frequently involving multiple components of the engine
- Also more recently: bugs in the parser!




Exploiting the Math.expm1 typing bug in V8

02 Jan 2019

Minus zero behaves like zero, right?



Optimization
Analysis
Other
“Breaks”




Type-Check Elimination

Type Inference

Runtim

CVE-2018-4233
CVE-2018-17463
CVE-2019-11707
CVE-2020-6418



(See “Attacking JavaScript Engines in 2022” by
@itszn13 and myself)

Write Barrier Elision

GC Modelling

CVE-2019-4442
CVE-2019-8622

Register Allocation

CVE-2018-12386

LICM
(loop-invariant code motion)

CVE-2019-8623
CVE-2019-8518

GVN
(global value numbering)

BCE
(bounds-check elimination)

Array Length Computation

Range Analysis

cbug 762874 (2017)
cbug 880207 (2018)
CVE-2019-13764

Pattern Matching

CVE-2021-30598
CVE-2021-30599


Lowering

CVE-2021-21220
cbug 1377775 (2022)

Spatial Memory Safety

Phase 3: Wasm Bugs

- Wasm first generally available in ~2017
- BUT: that was “linear memory” Wasm:
 - Simple instruction set (basically loads+stores, arithmetic, logic operations, jumps, calls)
 - No object model, only linear memory range (in which memory corruption didn’t matter)
 - Mostly only allowed compiling C/C++ to Wasm
- Then ~2023 WasmGC shipped
 - Major leap in complexity
 - Added powerful and highly complex object model + type system + garbage collection
 - ⇒ Major source of new vulnerabilities




A new way to bring garbage collected programming languages efficiently to WebAssembly

Published 01 November 2023 · Tagged with [WebAssembly](#)


WasmGC Bug Example (e.g. [339736513](#))

```
// ...
if (!IsJSReceiver(obj))
    return false;
if (IsJSProxy(obj))
    return false
Cast<JSObject>(obj)->Foo();
```



WasmGC Bug Example (e.g. [339736513](#))

```
// ...
if (!IsJSReceiver(obj))
    return false;
if (IsJSProxy(obj))
    return false
Cast<JSObject>(obj)->Foo();
```




Current and Future Defenses


Going Jitless?

Going Jitless?

- Idea: disable compilers
- Obvious attack surface reduction (always good!)
- In addition: may allow enabling additional mitigations (mostly W^X)
- But: potentially dramatic performance penalty (anything from 5%-90%)




Going Jitless? What about Wasm (in V8)?




- Problem: Wasm doesn't need/use an interpreter
- Options:
 - Add a Wasm interpreter
⇒ new attack surface : /
(unless written in memory-safe language?)
 - Disable Wasm : /

Going Jitless? What about Wasm (in V8)?



- Problem: Wasm doesn't need/use an interpreter
- Options:
 - Add a Wasm interpreter
⇒ new attack surface : /
(unless written in memory-safe language?)
 - Disable Wasm : /
 - Keep Baseline Wasm JIT : |
- Hypothesis: just disabling optimizing JITs is sufficient

Going Jitless? The Data!



- Tracking sheet for V8 bugs that are known to be exploitable
- Sources: ITW, exploit competitions, V8CTF
- Seems to confirm:
 - JITs account for ~50% of exploitable vulnerabilities
 - Baseline JITs are rarely the source of vulnerabilities

#	Issue	First Exploited	Description	Exploit requires V8 Sandbox Bypass	Exploit requires optimizing JITs (Turbofan & Maglev)	Exploit requires any JITs (Liftoff, Sparkplug, Maglev & Turbofan)	Variant	JavaScript or WebAssembly	Introduced by	Introduced in
1	386565144	V8CTF	Incorrect optimization in Maglev	Yes	Yes	Yes	No	JavaScript	Performance Work	2024
2	391907159	V8CTF	Wasm JIT allocation UaF	No*	Yes	Yes	No	WebAssembly	Performance Work	2024
3	398065918	V8CTF	Improper allocation folding in Maglev	Yes	Yes	Yes	No	JavaScript	Performance Work	2024
4	400052777	V8CTF	Incorrect handling of aliases during ElementsKind transitions	Yes	Yes	Yes	No	JavaScript	Performance Work	2024
5	403364367	V8CTF	Invalid handling of Wasm stack frames during stack walking	Yes	No	No	No	Both	Feature Work	2020
6	420636529	ITW	Logic error in store-store elimination	Yes	Yes	Yes	No	JavaScript	Performance Work	2024
7	422313191	TyphoonPWN	Invalid WebAssembly type canonicalization	Yes	No	No	Yes	WebAssembly	Performance Work	2025
8	427663123	ITW	Invalid hole check elision in interpreter	Yes	Probably	Probably	No	JavaScript	Performance Work	2023
9	433533359	V8CTF	Concurrent modification of Wasm code	No	No	No	No	N/A	Feature Work	2018 (?)
10	430344952	V8CTF	Divergence between preparsed and parser	Yes	No	No	No	JavaScript	Performance Work	2016 (?)
11	436181695	V8CTF	Invalid parsing of 'await using' in c-style loops	No	No	No	No	JavaScript	Feature Work	2025
12	445380761	ITW	Invalid integer optimization on Arm64	Yes	Yes	Yes	No	JavaScript	Performance Work	2016

Going Jitless?

Worth emphasizing:


50% fewer bugs != 50% fewer exploits

⇒ Jitless isn't the solution (but can still be a useful tool!)

Memory Tagging Extension (MTE)

Memory Tagging Extension (MTE)

- Hardware feature for memory safety enforcement
- Basic idea: add tag to pointers and memory, enforce that they match
- Can mitigate a number of bug types, e.g.
 - Linear OOB accesses
 - Use-after-Free
- See [ARM's documentation](#), [Project Zero's analysis](#), or [Apple's blog post](#) for details



<https://developer.arm.com/documentation/108035/0100/How-does-MTE-work>

Memory Tagging Extension (MTE)

- **Unlikely to have much impact in JavaScript engines**
 - (But most likely elsewhere!)
- Typical bugs there are too powerful:
 - Arbitrary OOB reads+writes
 - Arbitrary type confusions
- Also, custom pointer encodings leave no space for MTE tags...
- Apple seems to have come to a similar conclusion

Memory Integrity Enforcement vs. real-world exploit chains



🚫 Blocked by secure allocators

❗ Surviving step

* Blocked by EMTE

① Logical step

🚫* Blocked by secure allocators and EMTE

<https://security.apple.com/blog/memory-integrity-enforcement/>


Sandboxing

A different approach...

Idea:

- Accept that bugs will happen and that memory will be corrupted
- Limit which memory can be corrupted
- Make that a security boundary

=> Result: an in-process sandbox



Higher
Addresses

0xa48000000000

V8 Sandbox (1TB)

0xa38000000000

Lower
Addresses

Higher
Addresses

0xa48000000000

V8 Sandbox (1TB)

HeapObj1

HeapObj2


0xa38000000000

Lower
Addresses

Higher
Addresses


0xa48000000000


V8 Sandbox (1TB)





0xa38000000000

Lower
Addresses









Higher
Addresses

0xa48000000000

V8 Sandbox (1TB)

Basically: ban all raw pointers!

HeapObj1

HeapObj2

Offset from sandbox


Index

0xa38000000000

External Ptr Table	
0	Type + Pointer
1	Type + Pointer


External
Object

Lower
Addresses



Sandbox with Hardware Support?

- In the future, should be possible to “drop privileges” when executing JS or Wasm code
- Would be very similar to userspace/kernel split
- Built on top of current software-only sandbox
- Ideally: want to be able to run *untrusted machine code*



V8 Sandbox v2.0 - Basic Idea

Privileged Code

- Can read+write all memory
- Must be careful... (or memory safe)

Sandboxed Code

- Can only write inside the sandbox
- Ideally, bugs here don't matter


The diagram illustrates the V8 Sandbox v2.0 architecture. It consists of three nested rectangular regions. The innermost region is light blue and labeled "V8 Sandbox". The middle region is red and labeled "V8 Sandbox". The outermost region is also light blue and labeled "Privileged" Address Space twice, once at the top and once at the bottom.

“Privileged” Address Space

V8 Sandbox

“Privileged” Address Space

Privileged vs. Sandboxed Code



V8 Sandbox - “Progress Bar”

- Probably at ~v0.95 of the sandbox (software only)
- Most design-level issues taken care of, but implementation issues remain
- Already integrated into VRP ⇒ up to \$20k for high-quality bypasses
- Have a [prototype for hardware sandboxing](#) based on Intel PKEYs
 - But PKEYs are fairly limited, [ARM's upcoming POE2](#) is much more powerful!

Slightly annoying to bypass


Actually hard to bypass

Secure V8




JavaScript Engine with Sandbox

Software-based Sandbox



Hardware-based Sandbox



WebKit's JITCage

- Similar idea: sandbox untrusted code
- Requires special hardware features
- Different goal: prevents JIT-generated code from executing certain instructions (e.g. syscalls) or performing unsafe control-flow transfers
- Similar hardware capabilities available in ARM's upcoming POE2 



JITCage

- The following instructions can't be executed in the JITCage
 - RET
 - BR/BLR/BL
 - SVC
 - MRS/MSR

The Future of Sandboxing?

	JITCage	Software Sandbox	Hardware Sandbox
Applicable To	JIT Code (?)	JIT + Builtin Code	JIT + Builtin Code
Restricts Accessible Memory	No	Yes	Yes
Restricts Instruction Set	Yes	Mostly*	Yes
Restricts Control-Flow	Yes	Mostly*	Yes

* Requires additional code validation (think: [NaCL](#)) for actual security guarantees

Sandboxing - Summary

- Sandbox is a useful mitigation in itself
 - Attackers already now need multiple bugs (or one really good one) to exploit V8
- But it's also an ***architecture that enables powerful mitigations:***
 - Memory-safe languages, hardened C++, and/or MTE for privileged code
 - Hardware sandboxing or code validation in software for sandboxed code
- ⇒ Plausible path towards a secure, high-performance JS engine



The V8 Sandbox


Published 04 April 2024 · Tagged with [security](#)

After almost three years since the [initial design document](#) and [hundreds of CLs](#) in the meantime, the V8 Sandbox — a lightweight, in-process sandbox for V8 — has now progressed to the point where it is no

JS Vulnerability Research in 2025

Finding JS Engine Bugs in 2025

- The classics approaches are still going strong
- Manual code auditing
 - Time consuming, significant ramp-up work required
 - But allows searching for high-quality vulnerabilities
 - See e.g. some of the awesome bug reports by
[Seunghyun Lee \(@0x10n\)](#)
- Fuzzing
 - Can find really cool bugs, also less cool ones
 - Targeted fuzzing for e.g. variants or against specific components is still promising



Finding JS Engine Bugs in 2025

- But there are also new approaches:
- AI-powered vulnerability research
 - In particular [Google's Big Sleep system](#)
 - Combine latest AI reasoning models with purpose-built tools (code browser and debugger tool) and environment (JS shell)
 - Able to find bugs that fuzzers cannot (or at least struggle a lot with)

Bugs found so far this year by Big Sleep in V8



A screenshot of a search results page from Google's Issue Tracker. The search query in the bar is "componentid:1836411 title:V8". The results table has columns for Title, Assignee, Status, and ID. The table shows 11 rows of bugs, all assigned to "Assigned" and fixed, with IDs ranging from 457989902 to 436210783.

TITLE	ASSIGNEE	STATUS	ID
High impact issue in V8	Assigned	Assigned	457989902
High impact issue in V8	Assigned	Assigned	452319320
High impact issue in V8	Assigned	Assigned	450458037
High impact issue in V8	Assigned	Assigned	449910706
Medium impact issue in V8	Fixed	Fixed	444189152
High impact issue in V8: Integer truncation dur...	Fixed	Fixed	444141029
High impact issue in V8: Bytecode corruption ...	Fixed	Fixed	443875388
Low impact issue in V8	Fixed	Fixed	440296017
Medium impact issue in V8: ArrayBuffer Use-a...	Fixed	Fixed	439771269
High impact issue in V8: race condition during ...	Fixed	Fixed	439521654
High impact issue in V8: Bytecode corruption ...	Fixed	Fixed	436210783

Issue [436181695](#) (Google Big Sleep)

- Bug in bytecode compiler
- Mismatch in number of yield points between parser and compiler
- Leads to jump into the middle of bytecode
 - *Highly* exploitable and also bypasses the (current) sandbox
- Relatively easy to discover but fuzzers had no chance
 - They weren't yet aware of new syntax...

```
async function* bug() {
    for (await using x = { [Symbol.asyncDispose]() {} }; 1; ) {}
}
async function run() {
    for await (const x of bug()) {}
}
run();
```

Issue [443765373](#) (Google Big Sleep)

- Exception handler is encoded as offset in bytecode
- Offset is stored as 28 bit integer (256 MB)
- Max. bytecode size: 512 MB
- => Huge bytecode array will cause truncation of offset
- => Again arbitrary bytecode execution

```
function simple() {
    try {
        throw 42;
    } catch (e) {
        return e;
    }
}

[generated bytecode for function: simple]
0x2ad5001000e4 @ 0 : 1b ff f9 Mov <context>, r0
0x2ad5001000e7 @ 3 : 0d 2a LdaSmi [42]
0x2ad5001000e9 @ 5 : b5 Throw
0x2ad5001000ea @ 6 : d1 Star1
0x2ad5001000eb @ 7 : 8d f8 00 CreateCatchContext r1, [0]
0x2ad5001000ee @ 10 : d2 ...
...
0x2ad5001000f7 @ 19 : b7 Return
Handler Table (size = 16)
from to hdlr (prediction, data)
( 3, 6 ) -> 6 (prediction=1, data=0)
```


Issue [382005099](#) (V8 Team)

- Bad interaction between WebAudio and JavaScript engine
- WebAudio supports custom audio processing nodes defined in JS
- For performance reasons, WebAudio changes CPU's handling of floats
- ⇒ JavaScript code runs in unexpected CPU mode
- ⇒ Turns out this is exploitable
- Affected multiple browsers (e.g. also [CVE-2025-24213 in Safari](#))
- **Quite possibly my favorite bug of the last few years :)**

Background: Denormal Floats

- Floating point number: 1 sign bit, X mantissa bits, Y exponent bits
 - Final value: $(-1)^{\text{sign}} \times (1 + \text{mantissa}) \times 2^{\text{exponent} - \text{bias}}$
- Normal float numbers: no leading zero bits in mantissa
- Denormal float number: leading zero bits in mantissa
- For better performance: CPU can disable denormals \Rightarrow they become zero
- BUT: JavaScript spec assumes denormals are supported and \neq zero

In computer science, **subnormal numbers** are the subset of **denormalized numbers** (sometimes called **denormals**) that fill the **underflow gap** around zero in **floating-point arithmetic**. Any non-zero number with magnitude smaller than the smallest positive **normal number** is *subnormal*, while *denormal* can also refer to numbers outside that range.^[1]



An unaugmented floating-point system would contain only normalized numbers (indicated in red). Allowing denormalized numbers (blue) extends the system's range.

Issue 382005099 (trigger)

```
const denormal = 5E-324;

console.log(`Denormal float value outside processor: ${denormal}`);
// Prints "5e-324"

class DenormalDemoProcessor extends AudioWorkletProcessor {
process(inputs, outputs, parameters) {
  this.port.postMessage(`Denormal float value inside processor: ${denormal}`);
  // Prints "0" (!)
}
}

registerProcessor('denormal-demo-processor', DenormalDemoProcessor);
```

Issue 382005099 (PoC Exploit)

// See <https://crbug.com/382005099#comment19>

```
function poc(x) {
  let obj = {denormal: 5E-324};
  new Float64Array();

  let positive = (x & 1) + 1;
  let denormal = Math.min(obj.denormal, positive);
  let b = Object.is(denormal, 0);

  let n = b | 0;
  n *= 0xffffffff;
  let o = n + 1;

  let o_ = (Math.random() <= 1) ? o : undefined;
  let i = Math.sign(o_) * 64;
  let first = [1];
  first[i] = 2;

  let second = [1,2,3];
  return {first, second};
}
```

- Exploit abuses assumptions made by optimizing compiler
 - Essentially that $5E-324 \neq 0$
- Leads to incorrect range analysis
- Requires a few more tricks to make it work though
- Might also be exploitable in other ways, e.g. via bytecode compiler

Thank you!

Questions?

